

Bibliography

- [1] Abubakar Abid, Ali Abdalla, Ali Abid, Dawood Khan, Abdulrahman Alfozan, and James Zou. Gradio: Hassle-Free Sharing and Testing of ML Models in the Wild, 2019. URL <https://arxiv.org/abs/1906.02569>.
- [2] Samira Abnar and Willem Zuidema. Quantifying Attention Flow in Transformers. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2020. URL <https://arxiv.org/abs/2005.00928>.
- [3] Armen Aghajanyan, Sonal Gupta, and Luke Zettlemoyer. Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021. URL <https://arxiv.org/abs/2012.13255>.
- [4] Meta AI. The Llama 4 Herd: The Beginning of a New Era of Natively Multimodal AI. *Meta AI Blog*, 2025. URL <https://ai.meta.com/blog/llama-4-multimodal-intelligence/>.
- [5] Mistral AI. Mistral Large 2. *Mistral AI Blog*, 2024. URL <https://mistral.ai/news/mistral-large-2407/>.
- [6] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete Problems in AI Safety. *arXiv Preprint arXiv:1606.06565*, 2016. URL <https://arxiv.org/abs/1606.06565>.
- [7] Yanxu An, Li Shen, Yifan Xu, and Xinmei Liu. DPPO: Direct Divergence-Based Policy Optimization for Language Model Alignment. *arXiv Preprint arXiv:2503.14532*, 2025. URL <https://arxiv.org/abs/2503.14532>.
- [8] Marcin Andrychowicz, Filip Wolski, Alex Ray, et al. Hindsight Experience Replay. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [9] Anthropic. Model Context Protocol, 2024. URL <https://modelcontextprotocol.io>.
- [10] Anthropic. Building Effective Agents, 2024. URL <https://www.anthropic.com/research/building-effective-agents>.
- [11] Anthropic. The Claude 3 Model Family: Opus, Sonnet, Haiku. *Anthropic Technical Report*, 2024. URL <https://www.anthropic.com/news/claude-3-family>.
- [12] Anthropic. Natural Language Autoencoders: Interpreting Neural Networks with Natural Language Descriptions. *Anthropic Research Blog*, 2026. URL <https://www.anthropic.com/research/natural-language-autoencoders>.
- [13] Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to Retrieve, Generate, and Critique Through Self-Reflection. *arXiv Preprint arXiv:2310.11511*, 2023. URL <https://arxiv.org/abs/2310.11511>.
- [14] Mohammad Gheshlaghi Azar, Mark Rowland, Bilal Piot, et al. A General Theoretical Paradigm to Understand Learning from Human Feedback. *arXiv Preprint arXiv:2310.12036*, 2024. URL <https://arxiv.org/abs/2310.12036>.

- [15] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer Normalization. *arXiv Preprint arXiv:1607.06450*, 2016. URL <https://arxiv.org/abs/1607.06450>.
- [16] Yuntao Bai, Andy Jones, Kamal Ndousse, et al. Constitutional AI: Harmlessness from AI Feedback. *arXiv Preprint arXiv:2212.08073*, 2022. URL <https://arxiv.org/abs/2212.08073>.
- [17] Zhizhou Bai, Yixin Bao, et al. Kimi K2: Open Agentic Intelligence. *arXiv preprint arXiv:2507.20534*, 2025. URL <https://arxiv.org/abs/2507.20534>.
- [18] Payal Bajaj, Daniel Campos, Nick Craswell, et al. MS MARCO: A Human Generated Machine Reading COMprehension Dataset. *arXiv Preprint arXiv:1611.09268*, 2016.
- [19] Satandeep Banerjee and Alon Lavie. METEOR: An Automatic Metric for MT Evaluation with Improved Correlation with Human Judgments. In *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, 2005. URL <https://aclanthology.org/W05-0909/>.
- [20] Oren Barkan, Edan Hapon, Avi Caciularu, Ido Dagan, and Noam Koenigstein. Grad-SAM: Explaining Transformers via Gradient Self-Attention Maps. In *Proceedings of the 30th ACM International Conference on Information and Knowledge Management (CIKM)*, 2021. URL <https://arxiv.org/abs/2104.13299>.
- [21] Andrew G. Barto and Sridhar Mahadevan. Recent Advances in Hierarchical Reinforcement Learning. *Discrete Event Dynamic Systems*, 13:341–379, 2003.
- [22] Fabio Luigi Bellifemine, Giovanni Caire, and Dominic Greenwood. JADE: Developing Multi-Agent Systems with JADE, 2007.
- [23] Iz Beltagy, Matthew E. Peters, and Arman Cohan. Longformer: The Long-Document Transformer. *arXiv Preprint arXiv:2004.05150*, 2020. URL <https://arxiv.org/abs/2004.05150>.
- [24] Tim Berners-Lee, James Hendler, and Ora Lassila. The Semantic Web. *Scientific American*, 2001.
- [25] Maciej Besta, Nils Blach, Ales Kubicek, et al. Graph of Thoughts: Solving Elaborate Problems with Large Language Models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.
- [26] Luca Beurer-Kellner, Marc Fischer, and Martin Vechev. Prompting Is Programming: A Query Language for Large Language Models. In *Proceedings of PLDI*, 2023.
- [27] Rogerio Bonatti, Dan Zhao, Francesco Bonacci, et al. Windows Agent Arena: Evaluating Multi-Modal OS Agents at Scale. *arXiv Preprint arXiv:2409.08264*, 2024. URL <https://arxiv.org/abs/2409.08264>.
- [28] Edward de Bono. Six Thinking Hats, 1985.
- [29] Ralph Allan Bradley and Milton E. Terry. Rank Analysis of Incomplete Block Designs: I. The Method of Paired Comparisons. *Biometrika*, 1952. URL <https://www.jstor.org/stable/2334029>.
- [30] Trenton Bricken, Adly Templeton, Joshua Batson, et al. Towards Monosemanticity: Decomposing Language Models with Dictionary Learning. *Transformer Circuits Thread*, 2023. URL <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- [31] Greg Brockman, Vicki Cheung, Ludwig Pettersson, et al. OpenAI Gym, 2016.
- [32] Tom Brown, Benjamin Mann, Nick Ryder, et al. Language Models Are Few-Shot Learners. *NeurIPS*, 2020.

- [33] Chris Burges, Tal Shaked, Erin Renshaw, et al. Learning to Rank Using Gradient Descent. In *Proceedings of the 22nd International Conference on Machine Learning (ICML)*, 2005.
- [34] Christopher J. C. Burges, Robert Ragno, and Quoc V. Le. Learning to Rank with Nonsmooth Cost Functions. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2006.
- [35] Tianle Cai, Yuhong Li, Zhengyang Geng, et al. Medusa: Simple LLM Inference Acceleration Framework with Multiple Decoding Heads. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2401.10774>.
- [36] Brian Campbell, John Bradley, Nat Sakimura, and Torsten Lodderstedt. OAuth 2.0 Mutual-TLS Client Authentication and Certificate-Bound Access Tokens, 2020. URL <https://datatracker.ietf.org/doc/html/rfc8705>.
- [37] Murray Campbell, A. Joseph Hoane Jr., and Feng-hsiung Hsu. Deep Blue. *Artificial Intelligence*, 2002.
- [38] Zhe Cao, Tao Qin, Tie-Yan Liu, Ming-Feng Tsai, and Hang Li. Learning to Rank: From Pairwise Approach to Listwise Approach. In *Proceedings of the 24th International Conference on Machine Learning (ICML)*, 2007.
- [39] Nicholas Carlini, Florian Tramer, Eric Wallace, et al. Extracting Training Data from Large Language Models. *USENIX Security Symposium*, 2021. URL <https://arxiv.org/abs/2012.07805>.
- [40] Chainlit. Chainlit: Build Production-Ready Conversational AI Applications, 2024. URL <https://chainlit.io>.
- [41] Harrison Chase. LangChain, 2022. URL <https://github.com/langchain-ai/langchain>.
- [42] Kushal Chawla, Jaysa Ramirez, Rene Clever, Gale Lucas, Jonathan May, and Jonathan Gratch. CaSiNo: A Corpus of Campsite Negotiation Dialogues for Automatic Negotiation Systems. In *Proceedings of NAACL*, 2021.
- [43] Huayu Chen, Guande Zheng, Yimeng Kim, and Yiwen Chen. Noise Contrastive Alignment of Language Models with Explicit Rewards. *arXiv Preprint arXiv:2402.05369*, 2024. URL <https://arxiv.org/abs/2402.05369>.
- [44] Jianlv Chen, Shitao Luo, Mingzheng Zhang, Zheng Liu, Yingxia Xiao, and Defu Han. Hybrid Retrieval for Open-Domain Question Answering. *arXiv Preprint arXiv:2210.06029*, 2022. URL <https://arxiv.org/abs/2210.06029>.
- [45] Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. BGE M3-Embedding: Multi-Lingual, Multi-Functionality, Multi-Granularity Text Embeddings Through Self-Knowledge Distillation. *arXiv Preprint arXiv:2402.03216*, 2024.
- [46] Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance. *arXiv Preprint arXiv:2305.05176*, 2023. URL <https://arxiv.org/abs/2305.05176>.
- [47] Mark Chen, Jerry Tworek, Heewoo Jun, et al. Evaluating Large Language Models Trained on Code. *arXiv Preprint arXiv:2107.03374*, 2021. URL <https://arxiv.org/abs/2107.03374>.
- [48] Shouyuan Chen, Sherman Wong, Liangjian Chen, and Yuandong Tian. Extending Context Window of Large Language Models via Positional Interpolation. *arXiv Preprint arXiv:2306.15595*, 2023. URL <https://arxiv.org/abs/2306.15595>.
- [49] Tianqi Chen, Bing Xu, Chiyuan Zhang, and Carlos Guestrin. Training Deep Nets with Sublinear Memory Cost. *arXiv Preprint arXiv:1604.06174*, 2016.

- [50] Yuxin Chen et al. MOSAIC: Safe Agentic Training via Plan-Check-Act-Refuse. *arXiv preprint*, 2026.
- [51] Zhiyu Chen, Yiwei Deng, Ruiqi Zhang, Hao Sun, and Weizhu Chen. GSPO: Sequence-Level Policy Optimization for Language Model Alignment. *arXiv Preprint arXiv:2502.12459*, 2025. URL <https://arxiv.org/abs/2502.12459>.
- [52] Boris Cherny. Loops Are the Next Step from Agents. Talk at Meta @Scale Conference, 2026. Head of Claude Code, Anthropic.
- [53] Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. Mem0: Building Production-Ready AI Agents with Scalable Long-Term Memory. *arXiv Preprint arXiv:2504.19413*, 2025. URL <https://arxiv.org/abs/2504.19413>.
- [54] Seokhyun Choi, Hyunji Park, Doyoung Moon, Kyuyoung Kim, and Edward Choi. GOPO: Group Ordinal Policy Optimization for LLM Alignment with Non-Verifiable Rewards. *arXiv Preprint arXiv:2505.12948*, 2025. URL <https://arxiv.org/abs/2505.12948>.
- [55] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, et al. PaLM: Scaling Language Modeling with Pathways. *arXiv Preprint arXiv:2204.02311*, 2022.
- [56] Sayak Ray Chowdhury, Anush Chakraborty, Srinadh Natarajan, Alekh Agarwal, and David Sontag. Provably Robust DPO: Aligning Language Models with Noisy Feedback. *arXiv Preprint arXiv:2403.00409*, 2024. URL <https://arxiv.org/abs/2403.00409>.
- [57] Paul F. Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep Reinforcement Learning from Human Preferences. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. URL <https://arxiv.org/abs/1706.03741>.
- [58] Yung-Sung Chuang, Yujia Xie, Hongyin Luo, Yoon Kim, James Glass, and Pengcheng He. DoLA: Decoding by Contrasting Layers Improves Factuality in Large Language Models. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2309.03883>.
- [59] Hyung Won Chung, Le Hou, Shayne Longpre, et al. Scaling Instruction-Finetuned Language Models. *Journal of Machine Learning Research*, 2024. URL <https://arxiv.org/abs/2210.11416>.
- [60] Jacob Cohen. A Coefficient of Agreement for Nominal Scales. *Educational and Psychological Measurement*, 1960.
- [61] World Wide Web Consortium. Decentralized Identifiers (DIDs) v1.0, 2022. URL <https://www.w3.org/TR/did-core/>.
- [62] Gordon V. Cormack, Charles L. A. Clarke, and Stefan Buettcher. Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods. In *Proceedings of the 32nd International ACM SIGIR Conference*, 2009.
- [63] Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse Autoencoders Find Highly Interpretable Features in Language Models. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2309.08600>.
- [64] Tri Dao. FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2307.08691>.
- [65] Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2205.14135>.

- [66] Abhishek Das, Théophile Gerber, Georgia Gkioxari, Stefan Lee, Devi Parikh, and Dhruv Batra. TarMAC: Targeted Multi-Agent Communication. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 2019. URL <https://arxiv.org/abs/1810.11187>.
- [67] Databricks. DBRX: A New State-of-the-Art Open LLM. *Databricks Blog*, 2024. URL <https://www.databricks.com/blog/introducing-dbrx-new-state-art-open-llm>.
- [68] Google DeepMind. AlphaProof and AlphaGeometry 2: Solving Olympiad Geometry Without Human Demonstrations, 2024. URL <https://deepmind.google/discover/blog/ai-solves-imo-problems-at-silver-medal-level/>.
- [69] DeepSeek-AI. DeepSeek-V2: A Strong, Economical, and Efficient Mixture-of-Experts Language Model. *arXiv Preprint arXiv:2405.04434*, 2024. URL <https://arxiv.org/abs/2405.04434>.
- [70] DeepSeek-AI. DeepSeek-V3 Technical Report. *arXiv Preprint arXiv:2412.19437*, 2024. URL <https://arxiv.org/abs/2412.19437>.
- [71] DeepSeek-AI. DeepSeek-V4: Towards Highly Efficient Million-Token Context Intelligence. *arXiv preprint arXiv:2606.19348*, 2026. URL <https://arxiv.org/abs/2606.19348>.
- [72] Daya Guo DeepSeek-AI, Dejian Yang, et al. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. *arXiv Preprint arXiv:2501.12948*, 2025. URL <https://arxiv.org/abs/2501.12948>.
- [73] Xiang Deng, Yu Gu, Boyuan Zheng, et al. Mind2Web: Towards a Generalist Agent for the Web. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2306.06070>.
- [74] Michael Dennis, Natasha Jaques, Eugene Vinitzky, et al. Emergent Complexity and Zero-Shot Transfer via Unsupervised Environment Design. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [75] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient Finetuning of Quantized Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2305.14314>.
- [76] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of NAACL-HLT*, 2019. URL <https://arxiv.org/abs/1810.04805>.
- [77] Yixin Dong, Charlie Moon, Yuekai Wang, et al. XGrammar: Flexible and Efficient Structured Generation Engine for Large Language Models. *arXiv Preprint arXiv:2411.15100*, 2024.
- [78] Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. Improving Factuality and Reasoning in Language Models Through Multiagent Debate. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2305.14325>.
- [79] Yan Duan, John Schulman, Xi Chen, Peter L. Bartlett, Ilya Sutskever, and Pieter Abbeel. RL2: Fast Reinforcement Learning via Slow Reinforcement Learning. *arXiv Preprint arXiv:1611.02779*, 2016.
- [80] Edmund H. Durfee, Victor R. Lesser, and Daniel D. Corkill. Trends in Cooperative Distributed Problem Solving. *IEEE Transactions on Knowledge and Data Engineering*, 1989. URL <https://doi.org/10.1109/69.43404>.
- [81] Hermann Ebbinghaus. Über Das Gedächtnis: Untersuchungen Zur Experimentellen Psychologie, 1885.

- [82] Darren Edge, Ha Trinh, Newman Cheng, et al. From Local to Global: A Graph RAG Approach to Query-Focused Summarization. *arXiv Preprint arXiv:2404.16130*, 2024. URL <https://arxiv.org/abs/2404.16130>.
- [83] Vage Egiazarian, Andrei Panferov, Denis Kuznedelev, Elias Frantar, Artem Babber, and Dan Alistarh. Extreme Compression of Large Language Models via Additive Quantization. *arXiv Preprint arXiv:2401.06118*, 2024. URL <https://arxiv.org/abs/2401.06118>.
- [84] Arpad E Elo. The Rating of Chess Players, Past and Present, 1978.
- [85] Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. RAGAs: Automated Evaluation of Retrieval Augmented Generation. *arXiv Preprint arXiv:2309.15217*, 2023. URL <https://arxiv.org/abs/2309.15217>.
- [86] Kawin Ethayarajh. How Contextual Are Contextualized Word Representations? Comparing the Geometry of BERT, ELMo, and GPT-2 Embeddings. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019. URL <https://arxiv.org/abs/1909.00512>.
- [87] Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. KTO: Model Alignment as Prospect Theoretic Optimization. *arXiv Preprint arXiv:2402.01306*, 2024. URL <https://arxiv.org/abs/2402.01306>.
- [88] Hugging Face. OpenEnv: An Interface Library for RL Post-Training with Environments, 2025. URL <https://github.com/huggingface/OpenEnv>.
- [89] William Fedus, Barret Zoph, and Noam Shazeer. Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity. *Journal of Machine Learning Research*, 2022.
- [90] David Ferrucci, Eric Brown, Jennifer Chu-Carroll, et al. Building Watson: An Overview of the DeepQA Project. *AI Magazine*, 2010.
- [91] John Rupert Firth. A Synopsis of Linguistic Theory, 1930–1955. *Studies in Linguistic Analysis*, 1957.
- [92] Joseph L Fleiss. Measuring Nominal Scale Agreement Among Many Raters. *Psychological Bulletin*, 1971.
- [93] Foundation for Intelligent Physical Agents. FIPA ACL Message Structure Specification, 2002. URL <http://www.fipa.org/specs/fipa00061/>.
- [94] Thibault Formal, Carlos Lassance, Benjamin Piwowarski, and Stéphane Clinchant. SPLADE v2: Sparse Lexical and Expansion Model for Information Retrieval. *arXiv Preprint arXiv:2109.10086*, 2021.
- [95] Thibault Formal, Benjamin Piwowarski, and Stéphane Clinchant. SPLADE: Sparse Lexical and Expansion Model for First Stage Ranking. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2021.
- [96] Tiago Forte. Building a Second Brain: A Proven Method to Organize Your Digital Life and Unlock Your Creative Potential, 2022.
- [97] Nadime Francis, Alastair Green, Paolo Guagliardo, et al. Cypher: An Evolving Query Language for Property Graphs. In *Proceedings of the 2018 International Conference on Management of Data (SIGMOD)*, 2018.
- [98] Elias Frantar and Dan Alistarh. SparseGPT: Massive Language Models Can Be Accurately Pruned in One-Shot. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2301.00774>.

- [99] Elias Frantar, Saleh Ashkboos, Torsten Hoefler, and Dan Alistarh. GPTQ: Accurate Post-Training Quantization for Generative Pre-Trained Transformers. *arXiv Preprint arXiv:2210.17323*, 2023. URL <https://arxiv.org/abs/2210.17323>.
- [100] Yao Fu, Rameswar Panda, Xinyao Niu, et al. Data Engineering for Scaling Language Models to 128K Context. *arXiv Preprint arXiv:2402.10171*, 2024. URL <https://arxiv.org/abs/2402.10171>.
- [101] Yichao Fu, Peter Bailis, Ion Stoica, and Hao Zhang. Break the Sequential Dependency of LLM Inference Using Lookahead Decoding. *arXiv Preprint arXiv:2402.02057*, 2024. URL <https://arxiv.org/abs/2402.02057>.
- [102] Paul Furgale, Severin Klingler, James Nolan, Matt Staats, Gaia Di Lorenzo, Elisa Martinez Abad, Christian Schueler, Razvan Dinu, Alessio Devoto, Pascal Berard, Gal Kaplun, Elad Sarafian, Riccardo Roveri, Leon Derczynski, and Ricardo Silveira Cabral. NVIDIA-labs OO Agents: Native Python Object-Oriented Agents. Technical report, 2026. URL <https://arxiv.org/abs/2607.20709>.
- [103] Avigdor Gal, Ateret Anaby-Tavor, Alberto Trombetta, and Danilo Montesi. A Framework for Modeling and Evaluating Automatic Semantic Reconciliation. In *Proceedings of the 31st International Conference on Very Large Data Bases (VLDB)*, 2005. URL https://link.springer.com/chapter/10.1007/11896548_42.
- [104] Isabel O. Gallegos, Ryan A. Rossi, Joe Barber, Shanshan Tong, et al. Bias and Fairness in Large Language Models: A Survey. *Computational Linguistics*, 2024. URL <https://arxiv.org/abs/2309.00770>.
- [105] Leo Gao, John Schulman, and Jacob Hilton. Scaling Laws for Reward Model Overoptimization. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023.
- [106] Luyu Gao, Xueguang Ma, Jimmy Lin, and Jamie Callan. Precise Zero-Shot Dense Retrieval Without Relevance Labels. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023. URL <https://arxiv.org/abs/2212.10496>.
- [107] Javier García and Fernando Fernández. A Comprehensive Survey on Safe Reinforcement Learning. *Journal of Machine Learning Research*, 2015.
- [108] Google Gemini Team. Gemini 1.5: Unlocking Multimodal Understanding Across Millions of Tokens of Context. *arXiv Preprint arXiv:2403.05530*, 2024. URL <https://arxiv.org/abs/2403.05530>.
- [109] Georgi Gerganov. GGUF: GPT-Generated Unified Format (llama.cpp), 2023. URL <https://github.com/ggerganov/llama.cpp>.
- [110] Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer Feed-Forward Layers Are Key-Value Memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2021.
- [111] Fabian Gloeckle, Badr Youbi Idrissi, Baptiste Rozière, David Lopez-Paz, and Gabriel Synnaeve. Better & Faster Large Language Models via Multi-Token Prediction. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2404.19737>.
- [112] Charles A. E Goodhart. Problems of Monetary Management: The U.K. Experience. *Monetary Theory and Practice*, 1984.
- [113] Google. Agent2Agent (A2A) Protocol, 2025. URL <https://developers.google.com/agent2agent>.

- [114] Google DeepMind. Decoupled DiLoCo: Distributed Training Across Geographic Regions. arXiv preprint, 2026.
- [115] Google Research. Thinking to Recall: How Reasoning Unlocks Parametric Knowledge in LLMs. COLM 2026, 2026.
- [116] Alexander Gorbatenko. Online DPO with Synchronised Reference Model Updates. *TRL Documentation*, 2024.
- [117] Pierre-Paul Grassé. La Reconstruction Du Nid Et Les Coordinations Interindividuelles Chez Bellicositermes Natalensis Et Cubitermes Sp. La Théorie de La Stigmergie: Essai d’interprétation Du Comportement Des Termites Constructeurs. *Insectes Sociaux*, 1959.
- [118] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The Llama 3 Herd of Models. *arXiv Preprint arXiv:2407.21783*, 2024. URL <https://arxiv.org/abs/2407.21783>.
- [119] Alex Graves, Greg Wayne, Malcolm Reynolds, et al. Hybrid Computing Using a Neural Network with Dynamic External Memory. *Nature*, 2016.
- [120] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not What You’ve Signed up for: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISec)*, 2023.
- [121] JSON-RPC Working Group. JSON-RPC 2.0 Specification, 2010. URL <https://www.jsonrpc.org/specification>.
- [122] Albert Gu and Tri Dao. Mamba: Linear-Time Sequence Modeling with Selective State Spaces. *arXiv Preprint arXiv:2312.00752*, 2023. URL <https://arxiv.org/abs/2312.00752>.
- [123] Albert Gu and Tri Dao. Mamba: Linear-Time Sequence Modeling with Selective State Spaces. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2312.00752>.
- [124] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On Calibration of Modern Neural Networks. In *International Conference on Machine Learning (ICML)*, 2017.
- [125] Mandy Guo, Joshua Ainslie, David Uthus, et al. LongT5: Efficient Text-to-Text Transformer for Long Sequences. *Findings of the Association for Computational Linguistics: NAACL 2022*, 2022. URL <https://arxiv.org/abs/2112.07916>.
- [126] Ruiqi Guo, Philip Sun, Erik Lindgren, et al. Accelerating Large-Scale Inference with Anisotropic Vector Quantization. In *International Conference on Machine Learning (ICML)*, 2020.
- [127] Shangmin Guo, Biao Zhang, Tianlin Liu, et al. Direct Language Model Alignment from Online AI Feedback. *arXiv Preprint arXiv:2402.04792*, 2024. URL <https://arxiv.org/abs/2402.04792>.
- [128] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. REALM: Retrieval-Augmented Language Model Pre-Training. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020. URL <https://arxiv.org/abs/2002.08909>.
- [129] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2018. URL <https://arxiv.org/abs/1801.01290>.
- [130] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. Dream to Control: Learning Behaviors by Latent Imagination. In *Proceedings of the 8th International Conference on Learning Representations (ICLR)*, 2020. URL <https://arxiv.org/abs/1912.01603>.

- [131] Chi Han, Minlie Li, and Wenting Chen. SAPO: Soft Adaptive Policy Optimization for Efficient LLM Alignment. *arXiv Preprint arXiv:2503.01739*, 2025. URL <https://arxiv.org/abs/2503.01739>.
- [132] Daniel Han and Michael Han. Unsloth: Efficient LLM Fine-Tuning, 2024. URL <https://github.com/unslothai/unsloth>.
- [133] Dick Hardt. The OAuth 2.0 Authorization Framework, 2012. URL <https://datatracker.ietf.org/doc/html/rfc6749>.
- [134] Steve Harris and Andy Seaborne. SPARQL 1.1 Query Language, 2013. URL <https://www.w3.org/TR/sparql11-query/>.
- [135] Barbara Hayes-Roth. A Blackboard Architecture for Control. *Artificial Intelligence*, 1985. URL [https://doi.org/10.1016/0004-3702\(85\)90063-3](https://doi.org/10.1016/0004-3702(85)90063-3).
- [136] Soufiane Hayou, Nikhil Ghosh, and Bin Yu. LoRA+: Efficient Low Rank Adaptation of Large Models. *arXiv Preprint arXiv:2402.12354*, 2024. URL <https://arxiv.org/abs/2402.12354>.
- [137] Dan Hendrycks, Collin Burns, Steven Basart, et al. Measuring Massive Multitask Language Understanding. In *International Conference on Learning Representations (ICLR)*, 2021.
- [138] Ralf Herbrich, Tom Minka, and Thore Graepel. TrueSkill™: A Bayesian Skill Rating System. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2006. URL <https://proceedings.neurips.cc/paper/2006/hash/f44ee263952e65b3610b8ba51229d1f9-Abstract.html>.
- [139] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the Knowledge in a Neural Network. *arXiv Preprint arXiv:1503.02531*, 2015.
- [140] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, et al. Training Compute-Optimal Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2203.15556>.
- [141] Jiwoo Hong, Noah Lee, and James Thorne. ORPO: Monolithic Preference Optimization Without Reference Model. *arXiv Preprint arXiv:2403.07691*, 2024. URL <https://arxiv.org/abs/2403.07691>.
- [142] Sirui Hong, Mingchen Zhuge, Jonathan Chen, et al. MetaGPT: Meta Programming for a Multi-Agent Collaborative Framework, 2024.
- [143] Arian Hosseini, Xingdi Yuan, Pascal Poupart, Adam Trischler, and Yoshua Bengio. V-STaR: Training Verifiers for Self-Taught Reasoners. *arXiv Preprint arXiv:2402.06457*, 2024.
- [144] Neil Houlsby, Andrei Giber, Stanislaw Jastrzebski, et al. Parameter-Efficient Transfer Learning for NLP. In *International Conference on Machine Learning (ICML)*, 2019. URL <https://arxiv.org/abs/1902.00751>.
- [145] Pin-Lun Hsu, Yun Dai, Vignesh Kothapalli, et al. Liger-Kernel: Efficient Triton Kernels for LLM Training. *arXiv Preprint arXiv:2410.10989*, 2024. URL <https://arxiv.org/abs/2410.10989>.
- [146] Edward J. Hu, Yelong Shen, Phillip Wallis, et al. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations (ICLR)*, 2022. URL <https://arxiv.org/abs/2106.09685>.
- [147] Edward J. Hu, Yelong Shen, Phillip Wallis, et al. LoRA: Low-Rank Adaptation of Large Language Models. *arXiv Preprint arXiv:2106.09685*, 2022.
- [148] Jian Hu, Xibin Tao, Weixun Zhu, Sicheng Yang, Jingwen Liu, and Zilin Li. OpenRLHF: An Easy-to-Use, Scalable and High-Performance RLHF Framework. *arXiv Preprint arXiv:2405.11143*, 2024.

- [149] Qinghao Hu, Shang Yang, et al. Taming the Long-Tail: Efficient Reasoning RL Training with Adaptive Drafter. *arXiv Preprint arXiv:2511.16665*, 2025. URL <https://arxiv.org/abs/2511.16665>.
- [150] Shengding Hu, Yuge Tu, Xu Han, et al. MiniCPM: Unveiling the Potential of Small Language Models with Scalable Training Strategies. *arXiv Preprint arXiv:2404.06395*, 2024. URL <https://arxiv.org/abs/2404.06395>.
- [151] Qian Huang, Jian Vora, Percy Liang, and Jure Leskovec. MAgentBench: Evaluating Language Agents on Machine Learning Experimentation. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2310.03302>.
- [152] Yanping Huang, Youlong Cheng, Ankur Bapna, et al. GPipe: Efficient Training of Giant Neural Networks Using Pipeline Parallelism. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [153] Evan Hubinger, Carson Denison, Jesse Mu, et al. Sleeper Agents: Training Deceptive LLMs That Persist Through Safety Training. *arXiv Preprint arXiv:2401.05566*, 2024.
- [154] LangChain Inc. LangGraph: Build Stateful Multi-Actor Applications with LLMs, 2024. URL <https://github.com/langchain-ai/langgraph>.
- [155] LangChain Inc. LangGraph Studio: The First Agent IDE, 2024. URL <https://github.com/langchain-ai/langgraph-studio>.
- [156] Streamlit Inc. Streamlit: The Fastest Way to Build and Share Data Apps, 2024. URL <https://streamlit.io>.
- [157] Geoffrey Irving, Paul Christiano, and Dario Amodei. AI Safety via Debate. *arXiv Preprint arXiv:1805.00899*, 2018. URL <https://arxiv.org/abs/1805.00899>.
- [158] Max Jaderberg, Wojciech M. Czarnecki, Iain Dunning, et al. Human-Level Performance in 3D Multiplayer Games with Population-Based Reinforcement Learning. *Science*, 2019. URL <https://www.science.org/doi/10.1126/science.aau6249>.
- [159] Neel Jain, Ping-yeh Chiang, Yuxin Wen, et al. NEFTune: Noisy Embeddings Improve Instruction Finetuning. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2310.05914>.
- [160] Sarthak Jain and Byron C. Wallace. Attention Is Not Explanation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2019. URL <https://arxiv.org/abs/1902.10186>.
- [161] Eric Jang, Shixiang Gu, and Ben Poole. Categorical Reparameterization with Gumbel-Softmax. In *International Conference on Learning Representations (ICLR)*, 2017.
- [162] Soyeong Jeong, Jinheon Baek, Sukmin Cho, Sung Ju Hwang, and Jong C. Park. Adaptive-RAG: Learning to Adapt Retrieval-Augmented Large Language Models Through Question Complexity. *arXiv Preprint arXiv:2403.14403*, 2024. URL <https://arxiv.org/abs/2403.14403>.
- [163] Jiatao Ji, Adam Fisch, Jason Weston, and Sainbayar Sukhbaatar. Towards Exact Optimization of Language Model Alignment. *arXiv Preprint arXiv:2402.05369*, 2024. URL <https://arxiv.org/abs/2402.05369>.
- [164] Ziwei Ji, Nayeon Lee, Rita Frieske, et al. Survey of Hallucination in Natural Language Generation. *ACM Computing Surveys*, 2023. URL <https://arxiv.org/abs/2202.03629>.
- [165] Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, et al. Mistral 7B. *arXiv Preprint arXiv:2310.06825*, 2023. URL <https://arxiv.org/abs/2310.06825>.

- [166] Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, et al. Mixtral of Experts. *arXiv Preprint arXiv:2401.04088*, 2024. URL <https://arxiv.org/abs/2401.04088>.
- [167] Minqi Jiang, Edward Grefenstette, and Tim Rocktäschel. Prioritized Level Replay. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021. URL <https://arxiv.org/abs/2010.03934>.
- [168] Yibo Jiang et al. Reasoning Beyond Chain-of-Thought: A Latent Computational View. *arXiv preprint arXiv:2601.08058*, 2025. URL <https://arxiv.org/abs/2601.08058>.
- [169] Carlos E. Jimenez, John Yang, Alexander Wettig, et al. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2310.06770>.
- [170] Bowen Jin, Hansi Zeng, et al. Search-R1: Training LLMs to Reason and Leverage Search Engines with Reinforcement Learning. *arXiv Preprint arXiv:2503.09516*, 2025.
- [171] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-Scale Similarity Search with GPUs. *IEEE Transactions on Big Data*, 2021.
- [172] Michael Jones, John Bradley, and Nat Sakimura. JSON Web Token (JWT), 2015. URL <https://datatracker.ietf.org/doc/html/rfc7519>.
- [173] Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension. In *Proceedings of ACL*, 2017.
- [174] John Jumper, Richard Evans, Alexander Pritzel, et al. Highly Accurate Protein Structure Prediction with AlphaFold. *Nature*, 2021.
- [175] Saurav Kadavath, Tom Conerly, Amanda Askell, et al. Language Models (Mostly) Know What They Know. *arXiv Preprint arXiv:2207.05221*, 2022. URL <https://arxiv.org/abs/2207.05221>.
- [176] Damjan Kalajdzievski. A Rank Stabilization Scaling Factor for Fine-Tuning with LoRA. *arXiv Preprint arXiv:2312.03732*, 2023. URL <https://arxiv.org/abs/2312.03732>.
- [177] Jared Kaplan, Sam McCandlish, Tom Henighan, et al. Scaling Laws for Neural Language Models. *arXiv Preprint arXiv:2001.08361*, 2020. URL <https://arxiv.org/abs/2001.08361>.
- [178] Andrej Karpathy. Context Engineering for AI Agents. Social media post, 2025. June 2025.
- [179] Andrej Karpathy. AutoResearch: Automated ML Experimentation via Agent Loops. Open-source project, 2026. URL <https://github.com/karpathy/autoresearch>.
- [180] Vladimir Karpukhin, Barlas Oğuz, Sewon Min, et al. Dense Passage Retrieval for Open-Domain Question Answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020. URL <https://arxiv.org/abs/2004.04906>.
- [181] Junlong Ke, Zichen Wen, Weijia Li, Conghui He, and Linfeng Zhang. Respecting Self-Uncertainty in On-Policy Self-Distillation for Efficient LLM Reasoning. *arXiv preprint arXiv:2605.13255*, 2026. URL <https://arxiv.org/abs/2605.13255>.
- [182] Omar Khattab and Matei Zaharia. ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT. In *Proceedings of the 43rd International ACM SIGIR Conference*, 2020. URL <https://arxiv.org/abs/2004.12832>.
- [183] Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, et al. DSPy: Compiling Declarative Language Model Calls into Self-Improving Pipelines. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2310.03714>.

- [184] Seungone Kim, Se June Jang, et al. Distilling Step-by-Step! Outperforming Larger Language Models with Less Training Data and Smaller Model Sizes. *Findings of the ACL*, 2023. URL <https://arxiv.org/abs/2305.02301>.
- [185] Diederik P. Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization. In *International Conference on Learning Representations (ICLR)*, 2015. URL <https://arxiv.org/abs/1412.6980>.
- [186] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, et al. Overcoming Catastrophic Forgetting in Neural Networks. In *Proceedings of the National Academy of Sciences*, 2017.
- [187] Levente Kocsis and Csaba Szepesvári. Bandit Based Monte-Carlo Planning. In *European Conference on Machine Learning (ECML)*, 2006.
- [188] Jing Yu Koh, Robert Lo, Lawrence Jang, et al. VisualWebArena: Evaluating Multimodal Agents on Realistic Visual Web Tasks. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. URL <https://arxiv.org/abs/2401.13649>.
- [189] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large Language Models Are Zero-Shot Reasoners. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2205.11916>.
- [190] Dawid Jan Kopiczko, Tijmen Blankevoort, and Markus Nagel. VeRA: Vector-Based Random Matrix Adaptation. *arXiv Preprint arXiv:2310.11454*, 2024. URL <https://arxiv.org/abs/2310.11454>.
- [191] Vijay Anand Korthikanti, Jared Casper, Sangkug Lym, et al. Reducing Activation Recomputation in Large Transformer Models. In *Proceedings of Machine Learning and Systems (MLSys)*, 2023.
- [192] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. ImageNet Classification with Deep Convolutional Neural Networks. *NeurIPS*, 2012.
- [193] Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation. In *International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2302.09664>.
- [194] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, et al. Natural Questions: A Benchmark for Question Answering Research. *Transactions of the Association for Computational Linguistics*, 2019.
- [195] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, et al. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles (SOSP)*, 2023. URL <https://arxiv.org/abs/2309.06180>.
- [196] Heinrich Küttler, Nantas Nardelli, Alexander H. Miller, et al. The NetHack Learning Environment. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. URL <https://arxiv.org/abs/2006.13760>.
- [197] Timothée Lacroix, Guillaume Obozinski, and Nicolas Usunier. Tensor Decompositions for Temporal Knowledge Base Completion. In *International Conference on Learning Representations (ICLR)*, 2020. URL <https://arxiv.org/abs/2004.04926>.
- [198] Nathan Lambert, Jacob Morrison, Valentina Pyatkin, et al. Tulu 3: Pushing Frontiers in Open Language Model Post-Training. *arXiv Preprint arXiv:2411.15124*, 2024. URL <https://arxiv.org/abs/2411.15124>.

- [199] Angeliki Lazaridou and Marco Baroni. Emergent Multi-Agent Communication in the Deep Learning Era. *arXiv Preprint arXiv:2006.02419*, 2020. URL <https://arxiv.org/abs/2006.02419>.
- [200] Hung Le, Yue Wang, Akhilesh Deepak Gotmare, Silvio Savarese, and Steven Hoi. CodeRL: Mastering Code Generation Through Pretrained Models and Deep Reinforcement Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [201] Thibault Le Sellier De Chezelles, Maxime Gasse, Alexandre Drouin, Massimo Caccia, et al. The BrowserGym Ecosystem for Web Agent Research. *arXiv Preprint arXiv:2412.05467*, 2024. URL <https://arxiv.org/abs/2412.05467>.
- [202] Chankyu Lee, Rajarshi Roy, Menber Xu, et al. NV-Embed: Improved Techniques for Training LLMs as Generalist Embedding Models. *arXiv Preprint arXiv:2405.17428*, 2024.
- [203] Katherine Lee, Daphne Ippolito, Andrew Nystrom, et al. Deduplicating Training Data Makes Language Models Better. In *Proceedings of the 60th Annual Meeting of the ACL*, 2022. URL <https://arxiv.org/abs/2107.06499>.
- [204] Brian Lester, Rami Al-Rfou, and Noah Constant. The Power of Scale for Parameter-Efficient Prompt Tuning. In *Proceedings of the 2021 Conference on EMNLP*, 2021. URL <https://arxiv.org/abs/2104.08691>.
- [205] Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast Inference from Transformers via Speculative Decoding. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2211.17192>.
- [206] Mike Lewis, Denis Yarats, Yann N. Dauphin, Devi Parikh, and Dhruv Batra. Deal or No Deal? End-to-End Learning for Negotiation Dialogues. In *Proceedings of EMNLP*, 2017.
- [207] Mike Lewis, Yinhan Liu, Naman Goyal, et al. BART: Denoising Sequence-to-Sequence Pre-Training for Natural Language Generation, Translation, and Comprehension. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2020. URL <https://arxiv.org/abs/1910.13461>.
- [208] Patrick Lewis, Ethan Perez, Aleksandra Piktus, et al. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. URL <https://arxiv.org/abs/2005.11401>.
- [209] Shen Li, Yanli Zhao, Rohan Varma, et al. PyTorch Distributed: Experiences on Accelerating Data Parallel Training. In *Proceedings of the VLDB Endowment*, 2020.
- [210] Xian Li, Ari Holtzman, Daniel Fried, et al. Contrastive Decoding: Open-Ended Text Generation as Optimization. *ACL*, 2023.
- [211] Xiang Lisa Li and Percy Liang. Prefix-Tuning: Optimizing Continuous Prompts for Generation. In *Proceedings of the 59th Annual Meeting of the ACL*, 2021. URL <https://arxiv.org/abs/2101.00190>.
- [212] Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE: Speculative Sampling Requires Rethinking Feature Uncertainty. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2401.15077>.
- [213] Zehan Li, Xin Zhang, Yanzhao Zhang, Dingkun Long, Pengjun Xie, and Meishan Zhang. Towards General Text Embeddings with Multi-Stage Contrastive Learning. *arXiv Preprint arXiv:2308.03281*, 2023.
- [214] Zheng Li et al. UniClawBench: A Bilingual Benchmark for Interactive Agent Evaluation. *arXiv preprint*, 2026.

- [215] Hunter Lightman, Vineet Kosaraju, Yura Burda, et al. Let’s Verify Step by Step. *arXiv Preprint arXiv:2305.20050*, 2023.
- [216] Chin-Yew Lin. ROUGE: A Package for Automatic Evaluation of Summaries. In *Text Summarization Branches Out: Proceedings of the ACL-04 Workshop*, 2004. URL <https://aclanthology.org/W04-1013/>.
- [217] Ji Lin, Jiaming Tang, Haotian Tang, et al. AWQ: Activation-Aware Weight Quantization for LLM Compression and Acceleration. In *Proceedings of Machine Learning and Systems (MLSys)*, 2024. URL <https://arxiv.org/abs/2306.00978>.
- [218] Kevin Lin, Charlie Snell, Yu Wang, et al. Sleep-Time Compute: Beyond Inference Scaling at Test-Time. *arXiv Preprint arXiv:2504.13171*, 2025. URL <https://arxiv.org/abs/2504.13171>.
- [219] Long-Ji Lin. Self-Improving Reactive Agents Based on Reinforcement Learning, Planning and Teaching. *Machine Learning*, 1992.
- [220] Xiaoqiang Lin, Aritra Ghosh, Bryan Kian Hsiang Low, Anshumali Shrivastava, and Vijai Mohan. REFRAG: Rethinking RAG Based Decoding. *arXiv Preprint arXiv:2509.01092*, 2025.
- [221] Hao Liu, Matei Zaharia, and Pieter Abbeel. Ring Attention with Blockwise Transformers for Near-Infinite Context. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2310.01889>.
- [222] Haokun Liu, Derek Tam, Mohammed Muqeeth, et al. Few-Shot Parameter-Efficient Fine-Tuning Is Better and Cheaper Than in-Context Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2205.05638>.
- [223] Jiachang Liu, Dinghan Shen, Yizhe Zhang, Bill Dolan, Lawrence Carin, and Weizhu Chen. What Makes Good in-Context Examples for GPT-3? In *Proceedings of Deep Learning Inside Out (DeeLIO), ACL Workshop*, 2022. URL <https://arxiv.org/abs/2101.06804>.
- [224] Keller Jordan Liu, Yuchen Jin, Zhanhui Zhou, et al. Muon: An Optimizer for Hidden Layers in Neural Networks. *arXiv preprint*, 2025. URL <https://github.com/KellerJordan/Muon>.
- [225] Nelson F. Liu, Kevin Lin, John Hewitt, et al. Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics*, 2024. URL <https://arxiv.org/abs/2307.03172>.
- [226] Nelson F. Liu, Kevin Lin, John Hewitt, et al. Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics*, 2024. URL <https://arxiv.org/abs/2307.03172>.
- [227] Qiguang et al Liu. Answer Convergence as a Signal for Early Stopping in Reasoning. In *Proceedings of EMNLP*, 2025. URL <https://aclanthology.org/2025.emnlp-main.904/>.
- [228] Shih-Yang Liu, Chien-Yi Wang, Hongxu Yin, et al. DoRA: Weight-Decomposed Low-Rank Adaptation. *arXiv Preprint arXiv:2402.09353*, 2024. URL <https://arxiv.org/abs/2402.09353>.
- [229] Xiao Liu, Hao Yu, Hanchen Zhang, et al. AgentBench: Evaluating LLMs as Agents. *arXiv Preprint arXiv:2308.03688*, 2023.
- [230] Xingyu et al Liu. Learning Like Humans: Advancing LLM Reasoning Capabilities via Adaptive Difficulty Curriculum Learning. *arXiv Preprint arXiv:2505.08364*, 2025. URL <https://arxiv.org/abs/2505.08364>.

- [231] Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG Evaluation Using GPT-4 with Better Human Alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023. URL <https://arxiv.org/abs/2303.16634>.
- [232] Yihao Liu, Lefan Han, Yifan Tan, et al. Understanding and Mitigating the Pretraining Distribution Bias in GRPO. *arXiv Preprint arXiv:2505.07888*, 2025. URL <https://arxiv.org/abs/2505.07888>.
- [233] Yinhan Liu, Myle Ott, Naman Goyal, et al. RoBERTa: A Robustly Optimized BERT Pretraining Approach. *arXiv Preprint arXiv:1907.11692*, 2019. URL <https://arxiv.org/abs/1907.11692>.
- [234] Zechun Liu, Barlas Oguz, Changsheng Zhao, et al. LLM-QAT: Data-Free Quantization Aware Training for Large Language Models. *arXiv Preprint arXiv:2305.17888*, 2023. URL <https://arxiv.org/abs/2305.17888>.
- [235] Zhuokai Liu, Hao Dong, et al. DLER: Doing Length Penalty Right—Incentivizing More Intelligence Per Token. *arXiv Preprint arXiv:2510.15110*, 2025. URL <https://arxiv.org/abs/2510.15110>.
- [236] Zichen Liu et al. OctoThinker: Mid-Training Determines RL-Readiness of Base Models. *arXiv preprint*, 2025.
- [237] Zirui Liu, Jiayi Yuan, Hongye Jin, et al. KIVI: A Tuning-Free Asymmetric 2bit Quantization for KV Cache. In *International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2402.02750>.
- [238] Ilya Loshchilov and Frank Hutter. Decoupled Weight Decay Regularization. *arXiv Preprint arXiv:1711.05101*, 2019. URL <https://arxiv.org/abs/1711.05101>.
- [239] Ryan Lowe, Yi Wu, Aviv Tamar, Jean Harb, Pieter Abbeel, and Igor Mordatch. Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. URL <https://arxiv.org/abs/1706.02275>.
- [240] Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. *arXiv Preprint arXiv:2408.06292*, 2024. URL <https://arxiv.org/abs/2408.06292>.
- [241] Yao Lu, Max Bartolo, Alastair Moore, Sebastian Riedel, and Pontus Stenetorp. Fantastically Ordered Prompts and Where to Find Them: Overcoming Few-Shot Prompt Order Sensitivity. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2022. URL <https://arxiv.org/abs/2104.08786>.
- [242] Zhenyu Luo, Ziyang Chen, Yulun Jiang, et al. ScaleRL: Scaling Reinforcement Learning for LLM Reasoning. *arXiv Preprint arXiv:2505.16356*, 2025. URL <https://arxiv.org/abs/2505.16356>.
- [243] Zhixun Luo, Jiaqi Shi, Tao Yu, and Minlie Chen. VESPO: Variational Sequence-Level Soft Policy Optimization for LLM Alignment. *arXiv Preprint arXiv:2505.07508*, 2025. URL <https://arxiv.org/abs/2505.07508>.
- [244] Tobi Lütke. Context Engineering. Social media post, 2025. CEO of Shopify, June 2025.
- [245] Jakub Lála, Odhran O’Donoghue, Aleksandar Shtedritski, Sam Cox, Samuel G. Rodrigues, and Andrew D. White. PaperQA: Retrieval-Augmented Generative Agent for Scientific Research. *arXiv Preprint arXiv:2312.07559*, 2023. URL <https://arxiv.org/abs/2312.07559>.
- [246] Aman Madaan, Niket Tandon, Prakhar Gupta, et al. Self-Refine: Iterative Refinement with Self-Feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

- [247] Adyasha Maharana, Dong-Ho Lee, Sergey Tuber, Mohit Ruber, Francesco Barbieri, and Mohit Bansal. LoCoMo: Long-Context Conversation with Memory Operations. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.
- [248] Yu. A. Malkov and D. A. Yashunin. Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2020.
- [249] Potsawee Manakul, Adian Liusie, and Mark J. F. Gales. SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. In *Proceedings of EMNLP*, 2023. URL <https://arxiv.org/abs/2303.08896>.
- [250] Sam McCandlish, Jared Kaplan, Dario Amodei, and OpenAI Dota Team. An Empirical Model of Large-Batch Training. *arXiv Preprint arXiv:1812.06162*, 2018.
- [251] Jiangfei et al Mei. APRIL: Active Partial Rollouts in Reinforcement Learning to Tame Long-Tail Generation. *arXiv Preprint arXiv:2509.18521*, 2025. URL <https://arxiv.org/abs/2509.18521>.
- [252] Yu Meng, Mengzhou Xia, and Danqi Chen. SimPO: Simple Preference Optimization with a Reference-Free Reward. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL <https://arxiv.org/abs/2405.14734>.
- [253] Meta AI. Remember When It Matters: Proactive Memory for Long-Horizon Agents. arXiv preprint, 2026.
- [254] Grégoire Mialon, Clémentine Fourrier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. GAIA: A Benchmark for General AI Assistants. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2311.12983>.
- [255] Paulius Micikevicius, Sharan Narang, Jonah Alben, et al. Mixed Precision Training. In *International Conference on Learning Representations (ICLR)*, 2018.
- [256] Microsoft. Semantic Kernel: SDK for Integrating AI Models into Applications, 2023. URL <https://github.com/microsoft/semantic-kernel>.
- [257] Sewon Min, Xinxu Lyu, Ari Holtzman, et al. Rethinking the Role of Demonstrations: What Makes in-Context Learning Work? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022. URL <https://arxiv.org/abs/2202.12837>.
- [258] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, et al. Human-Level Control Through Deep Reinforcement Learning. *Nature*, 2015. URL <https://www.nature.com/articles/nature14236>.
- [259] Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, et al. Asynchronous Methods for Deep Reinforcement Learning. In *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, 2016. URL <https://arxiv.org/abs/1602.01783>.
- [260] MosaicML. MPT-7B: A New Standard for Open-Source, Commercially Usable LLMs. *MosaicML Blog*, 2023. URL <https://www.mosaicml.com/blog/mpt-7b>.
- [261] João Moura. CrewAI: Framework for Orchestrating Role-Playing Autonomous AI Agents, 2023. URL <https://github.com/crewAIInc/crewAI>.
- [262] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, et al. WebGPT: Browser-Assisted Question-Answering with Human Feedback. *arXiv Preprint arXiv:2112.09332*, 2021. URL <https://arxiv.org/abs/2112.09332>.
- [263] Deepak Narayanan, Aaron Harlap, Amar Phanishayee, et al. PipeDream: Generalized Pipeline Parallelism for DNN Training. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles (SOSP)*, 2019.

- [264] Deepak Narayanan, Mohammad Shoeybi, Jared Casper, et al. Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM. *arXiv Preprint arXiv:2104.04473*, 2021.
- [265] Andrew Y. Ng, Daishi Harada, and Stuart J. Russell. Policy Invariance Under Reward Transformations: Theory and Application to Reward Shaping. In *Proceedings of the 16th International Conference on Machine Learning (ICML)*, 1999.
- [266] Minh Nguyen. Min-p Sampling: A Simple Baseline for Better LLM Decoding. *arXiv Preprint arXiv:2310.06022*, 2024.
- [267] Noam Nisan, Tim Roughgarden, Éva Tardos, and Vijay V. Vazirani. Algorithmic Game Theory, 2007. URL <https://www.cs.cmu.edu/~sandholm/cs15-892F13/algorithmic-game-theory.pdf>.
- [268] Rodrigo Nogueira and Kyunghyun Cho. Passage Re-Ranking with BERT. *arXiv Preprint arXiv:1901.04085*, 2019.
- [269] Rodrigo Nogueira, Zhiying Jiang, Ronak Pradeep, and Jimmy Lin. Document Ranking with a Pretrained Sequence-to-Sequence Model. *Findings of EMNLP*, 2020.
- [270] NVIDIA. Dynamo: A Datacenter-Scale Distributed Inference Serving Framework. Open-source project, 2025. URL <https://github.com/ai-dynamo/dynamo>.
- [271] NVIDIA. NeMo Gym: Interactive RL Environments for LLM Agent Training. Technical report, 2025. URL <https://developer.nvidia.com/nemo-gym>.
- [272] Maxwell Nye, Anders Johan Andreassen, Guy Gur-Ari, et al. Show Your Work: Scratchpads for Intermediate Computation with Language Models. *arXiv Preprint arXiv:2112.00114*, 2021. URL <https://arxiv.org/abs/2112.00114>.
- [273] Catherine Olsson, Nelson Elhage, Neel Nanda, et al. In-Context Learning and Induction Heads. *Transformer Circuits Thread*, 2022. URL <https://transformer-circuits.pub/2022/in-context-learning-and-induction-heads/index.html>.
- [274] OpenAI. GPT-4 Technical Report. *arXiv Preprint arXiv:2303.08774*, 2023.
- [275] OpenAI. Learning to Reason with LLMs, 2024. URL <https://openai.com/index/learning-to-reason-with-llms/>.
- [276] OpenAI. Swarm: An Educational Framework for Lightweight Multi-Agent Orchestration, 2024. URL <https://github.com/openai/swarm>.
- [277] OpenAI. OpenAI Agents SDK, 2025. URL <https://github.com/openai/openai-agents-python>.
- [278] OpenAI. OpenAI o3 and o4-mini System Card, 2025. URL <https://openai.com/index/o3-and-o4-mini-system-card/>.
- [279] Addy Osmani. Loop Engineering. Blog post, 2026. URL <https://addyosmani.com/blog/loop-engineering/>.
- [280] Long Ouyang, Jeffrey Wu, Xu Jiang, et al. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2203.02155>.
- [281] Charles Packer, Vivian Fang, Shishir G. Patil, Kevin Lin, Sarah Wooders, and Joseph E. Gonzalez. MemGPT: Towards LLMs as Operating Systems. *arXiv Preprint arXiv:2310.08560*, 2023. URL <https://arxiv.org/abs/2310.08560>.

- [282] Leyi Pan, Shuchang Tao, Yunpeng Zhai, Lingzhe Zhang, Zhaoyang Liu, Bolin Ding, Aiwei Liu, and Lijie Wen. RLCSO: Reinforcement Learning with Contrastive On-Policy Self-Distillation. *arXiv preprint arXiv:2606.11709*, 2026. URL <https://arxiv.org/abs/2606.11709>.
- [283] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. BLEU: A Method for Automatic Evaluation of Machine Translation. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2002. URL <https://aclanthology.org/P02-1040/>.
- [284] Raja Parasuraman and Victor Riley. Humans and Automation: Use, Misuse, Disuse, Abuse. *Human Factors*, 1997. URL <https://doi.org/10.1518/001872097778543886>.
- [285] Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative Agents: Interactive Simulacra of Human Behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*, 2023. URL <https://arxiv.org/abs/2304.03442>.
- [286] Deepak Pathak, Pulkit Agrawal, Alexei A. Efros, and Trevor Darrell. Curiosity-Driven Exploration by Self-Supervised Prediction. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 2017.
- [287] Shishir G. Patil, Tianjun Zhang, Xin Wang, and Joseph E. Gonzalez. Gorilla: Large Language Model Connected with Massive APIs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL <https://arxiv.org/abs/2305.15334>.
- [288] Bo Peng, Eric Alcaide, Quentin Anthony, et al. RWKV: Reinventing RNNs for the Transformer Era. *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. URL <https://arxiv.org/abs/2305.13048>.
- [289] Bowen Peng, Jeffrey Quesnelle, Honglu Fan, and Enrico Shao. YaRN: Efficient Context Window Extension of Large Language Models. *arXiv Preprint arXiv:2309.00071*, 2023.
- [290] Ethan Perez, Saffron Huang, Francis Song, et al. Red Teaming Language Models with Language Models. In *Proceedings of EMNLP*, 2022. URL <https://arxiv.org/abs/2202.03286>.
- [291] R. L Plackett. The Analysis of Permutations. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 1975.
- [292] Ofir Press, Noah A. Smith, and Mike Lewis. Train Short, Test Long: Attention with Linear Biases Enables Input Length Generalization. *ICLR*, 2022.
- [293] Tom Preston-Werner. Semantic Versioning 2.0.0, 2024. URL <https://semver.org/>.
- [294] Pranav Putta, Edmund Mills, Naman Garg, et al. Agent Q: Advanced Reasoning and Learning for Autonomous AI Agents. *arXiv Preprint arXiv:2408.07199*, 2024.
- [295] PyTorch Team. Miles: A PyTorch-Native Stack for Large-Scale LLM RL Post-Training. PyTorch Blog, 2026. URL <https://pytorch.org/blog/miles-a-pytorch-native-stack-for-large-scale-llm-rl-post-training/>.
- [296] Penghui Qi, Xinyi Wan, Guangxing Huang, and Min Lin. Zero Bubble Pipeline Parallelism. *arXiv Preprint arXiv:2401.10241*, 2023.
- [297] Zhenting Qi, Mingyuan Wan, Jialin Cao, and Min Lin. Mutual Reasoning Makes Smaller LLMs Stronger Problem-Solvers. *arXiv Preprint arXiv:2408.06195*, 2024.
- [298] Yiwei Qin, Xuefeng Li, Haoyang Zou, et al. O1 Replication Journey: A Strategic Progress Report. *arXiv Preprint arXiv:2410.18982*, 2024. URL <https://arxiv.org/abs/2410.18982>.

- [299] Yujia Qin, Shihao Liang, Yining Ye, et al. ToolLLM: Facilitating Large Language Models to Master 16000+ Real-World APIs. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2307.16789>.
- [300] Adrian H Rackauckas. RAG-Fusion: A New Take on Retrieval-Augmented Generation, 2024.
- [301] Alec Radford, Jeffrey Wu, Rewon Child, David Luen, Dario Amodei, and Ilya Sutskever. Language Models Are Unsupervised Multitask Learners. *OpenAI Blog*, 2019. URL https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.
- [302] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. Direct Preference Optimization: Your Language Model Is Secretly a Reward Model. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2305.18290>.
- [303] Colin Raffel, Noam Shazeer, Adam Roberts, et al. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. *Journal of Machine Learning Research*, 2020. URL <https://arxiv.org/abs/1910.10683>.
- [304] Samyam Rajbhandari, Jeff Rasley, Olatunji Rber, and Yuxiong He. ZeRO: Memory Optimizations Toward Training Trillion Parameter Models. *arXiv Preprint arXiv:1910.02054*, 2020.
- [305] Samyam Rajbhandari, Olatunji Ruwase, Jeff Rasley, Shaden Smith, and Yuxiong He. ZeRO-Infinity: Breaking the GPU Memory Wall for Extreme Scale Deep Learning. *arXiv Preprint arXiv:2104.07857*, 2021.
- [306] Tabish Rashid, Mikayel Samvelyan, Christian Schroeder de Witt, Gregory Farquhar, Jakob Foerster, and Shimon Whiteson. QMIX: Monotonic Value Function Factorisation for Deep Multi-Agent Reinforcement Learning. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2018. URL <https://arxiv.org/abs/1803.11605>.
- [307] Eric Rescorla. The Transport Layer Security (TLS) Protocol Version 1.3, 2018. URL <https://datatracker.ietf.org/doc/html/rfc8446>.
- [308] Herbert Robbins and Sutton Monro. A Stochastic Approximation Method. *The Annals of Mathematical Statistics*, 1951.
- [309] Stephen Robertson and Hugo Zaragoza. The Probabilistic Relevance Framework: BM25 and Beyond. *Foundations and Trends in Information Retrieval*, 2009.
- [310] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. Learning Representations by Back-Propagating Errors. *Nature*, 1986. URL <https://doi.org/10.1038/323533a0>.
- [311] Gavin A. Rummery and Mahesan Niranjan. On-Line q-Learning Using Connectionist Systems, 1994.
- [312] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. DistilBERT, a Distilled Version of BERT: Smaller, Faster, Cheaper and Lighter. *arXiv Preprint arXiv:1910.01108*, 2019.
- [313] Keshav Santhanam, Omar Khattab, Jon Saad-Falcon, Christopher Potts, and Matei Zaharia. ColBERTv2: Effective and Efficient Retrieval via Lightweight Late Interaction. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2022.
- [314] Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. Prioritized Experience Replay. In *Proceedings of the 4th International Conference on Learning Representations (ICLR)*, 2016. URL <https://arxiv.org/abs/1511.05952>.

- [315] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, et al. Toolformer: Language Models Can Teach Themselves to Use Tools. In *Advances in Neural Information Processing Systems*, 2023.
- [316] Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, et al. Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model. *Nature*, 2020. URL <https://arxiv.org/abs/1911.08265>.
- [317] John Schulman, Sergey Levine, Pieter Abbeel, Michael Jordan, and Philipp Moritz. Trust Region Policy Optimization. In *Proceedings of the 32nd International Conference on Machine Learning (ICML)*, 2015. URL <https://arxiv.org/abs/1502.05477>.
- [318] John Schulman, Philipp Moritz, Sergey Levine, Michael Jordan, and Pieter Abbeel. High-Dimensional Continuous Control Using Generalized Advantage Estimation. In *Proceedings of the 4th International Conference on Learning Representations (ICLR)*, 2016. URL <https://arxiv.org/abs/1506.02438>.
- [319] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal Policy Optimization Algorithms. *arXiv Preprint arXiv:1707.06347*, 2017. URL <https://arxiv.org/abs/1707.06347>.
- [320] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural Machine Translation of Rare Words with Subword Units. In *Proceedings of the 54th Annual Meeting of the ACL*, 2016. URL <https://arxiv.org/abs/1508.07909>.
- [321] Alexander Sergeev and Mike Del Balso. Horovod: Fast and Easy Distributed Deep Learning in TensorFlow. *arXiv Preprint arXiv:1802.05799*, 2018.
- [322] Jay Shah, Ganesh Bikshandi, Ying Zhang, Vijay Thakkar, Pradeep Ramani, and Tri Dao. FlashAttention-3: Fast and Accurate Attention with Asynchrony and Low-Precision. *arXiv Preprint arXiv:2407.08691*, 2024. URL <https://arxiv.org/abs/2407.08691>.
- [323] Zhihong Shao, Peiyi Wang, Qihao Zhu, et al. DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models. *arXiv Preprint arXiv:2402.03300*, 2024. URL <https://arxiv.org/abs/2402.03300>.
- [324] Lloyd S Shapley. Stochastic Games. In *Proceedings of the National Academy of Sciences*, 1953. URL <https://www.pnas.org/doi/10.1073/pnas.39.10.1095>.
- [325] Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarsz, et al. Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer. In *International Conference on Learning Representations (ICLR)*, 2017. URL <https://arxiv.org/abs/1701.06538>.
- [326] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language Agents with Verbal Reinforcement Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2303.11366>.
- [327] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-LM: Training Multi-Billion Parameter Language Models Using Model Parallelism. *arXiv Preprint arXiv:1909.08053*, 2019.
- [328] Yoav Shoham and Kevin Leyton-Brown. Multiagent Systems: Algorithmic, Game-Theoretic, and Logical Foundations, 2008. URL <http://www.masfoundations.org/>.
- [329] Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. ALFWorld: Aligning Text and Embodied Environments for Interactive Learning. In *International Conference on Learning Representations (ICLR)*, 2021.
- [330] Significant Gravitas. AutoGPT: An Autonomous GPT-4 Experiment. Open-source project, 2023. URL <https://github.com/Significant-Gravitas/AutoGPT>.

- [331] David Silver, Aja Huang, Chris J. Maddison, et al. Mastering the Game of Go with Deep Neural Networks and Tree Search. *Nature*, 2016. URL <https://www.nature.com/articles/nature16961>.
- [332] David Silver, Julian Schrittwieser, Karen Simonyan, et al. Mastering the Game of Go Without Human Knowledge. *Nature*, 2017. URL <https://www.nature.com/articles/nature24270>.
- [333] Joar Skalse, Nikolaus Howe, Dmitrii Krashennnikov, and David Krueger. Defining and Characterizing Reward Hacking. In *Advances in Neural Information Processing Systems*, 2022. URL <https://arxiv.org/abs/2209.13085>.
- [334] Reid G Smith. The Contract Net Protocol: High-Level Communication and Control in a Distributed Problem Solver. *IEEE Transactions on Computers*, 1980. URL <https://doi.org/10.1109/TC.1980.1675516>.
- [335] Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM Test-Time Compute Optimally Can Be More Effective Than Scaling Model Parameters. *arXiv Preprint arXiv:2408.03314*, 2024.
- [336] Karen Sparck Jones. A Statistical Interpretation of Term Specificity and Its Application in Retrieval. *Journal of Documentation*, 1972.
- [337] Larry R Squire. Declarative and Nondeclarative Memory: Multiple Brain Systems Supporting Learning and Memory. *Journal of Cognitive Neuroscience*, 1992. URL <https://doi.org/10.1162/jocn.1992.4.3.232>.
- [338] Peter Steinberger. OpenClaw: An Always-On Agent Orchestrator. Open-source project, 2026. URL <https://github.com/openclaw>.
- [339] Nisan Stiennon, Long Ouyang, Jeffrey Wu, et al. Learning to Summarize from Human Feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. URL <https://arxiv.org/abs/2009.01325>.
- [340] Jianlin Su, Jiarun Cao, Weijie Liu, and Yangyiwen Ou. Whitening Sentence Representations for Better Semantics and Faster Retrieval. *arXiv Preprint arXiv:2103.15316*, 2021. URL <https://arxiv.org/abs/2103.15316>.
- [341] Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. RoFormer: Enhanced Transformer with Rotary Position Embedding. *Neurocomputing*, 2024.
- [342] Sainbayar Sukhbaatar, Arthur Szlam, Jason Weston, and Rob Fergus. End-to-End Memory Networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2015. URL <https://arxiv.org/abs/1503.08895>.
- [343] Sainbayar Sukhbaatar, Arthur Szlam, and Rob Fergus. Learning Multiagent Communication with Backpropagation. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2016. URL <https://arxiv.org/abs/1605.07736>.
- [344] Theodore R. Sumers, Shunyu Yao, Karthik Narasimhan, and Thomas L. Griffiths. Cognitive Architectures for Language Agents. *Transactions on Machine Learning Research (TMLR)*, 2024. URL <https://arxiv.org/abs/2309.02427>.
- [345] Mingjie Sun, Zhuang Liu, Anna Bair, and J. Zico Kolter. A Simple and Effective Pruning Approach for Large Language Models. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2306.11695>.
- [346] Richard S Sutton. Learning to Predict by the Methods of Temporal Differences. *Machine Learning*, 1988.

- [347] Richard S. Sutton and Andrew G. Barto. Reinforcement Learning: An Introduction, 2018. URL <http://incompleteideas.net/book/the-book-2nd.html>.
- [348] Meta PyTorch Team. TorchForge: PyTorch-Native Post-Training at Scale, 2025. URL <https://github.com/meta-pytorch/torchforge>.
- [349] PyTorch Team. TorchTune: PyTorch Native Post-Training Library, 2024. URL <https://github.com/pytorch/torchTune>.
- [350] Qwen Team. Qwen2.5: A Party of Foundation Models. *arXiv Preprint arXiv:2412.15115*, 2024. URL <https://arxiv.org/abs/2412.15115>.
- [351] Qwen Team. QwQ: Reflect Deeply on the Boundaries of the Unknown. *Qwen Blog*, 2024. URL <https://qwenlm.github.io/blog/qwq-32b-preview/>.
- [352] Qwen Team. Qwen3 Technical Report. *arXiv Preprint arXiv:2505.09388*, 2025.
- [353] Adly Templeton, Tom Conerly, Jonathan Marcus, et al. Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet. *Transformer Circuits Thread*, 2024. URL <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html>.
- [354] Eric Todd, Millicent L. Li, Arnab Sen Sharma, Aaron Mueller, Byron C. Wallace, and David Bau. Function Vectors in Large Language Models. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2310.15213>.
- [355] Mark Towers, Ariel Kwiatkowski, Jordan Terry, et al. Gymnasium: A Standard Interface for Reinforcement Learning Environments. *NeurIPS Datasets and Benchmarks*, 2024. URL <https://arxiv.org/abs/2407.17032>.
- [356] Endel Tulving. Memory and Consciousness. *Canadian Psychology / Psychologie Canadienne*, 1985.
- [357] Ashish Vaswani, Noam Shazeer, Niki Parmar, et al. Attention Is All You Need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. URL <https://arxiv.org/abs/1706.03762>.
- [358] Vercel. Vercel AI SDK, 2024. URL <https://sdk.vercel.ai>.
- [359] Jesse Vig. A Multiscale Visualization of Attention in the Transformer Model. In *Proceedings of the 57th ACL: System Demonstrations*, 2019. URL <https://arxiv.org/abs/1906.05714>.
- [360] Ashwin K. Vijayakumar, Michael Cogswell, Ramprasaath R. Selvaraju, et al. Diverse Beam Search: Decoding Diverse Solutions from Neural Sequence Models. *AAAI*, 2018.
- [361] Elena Voita, David Talbot, Fedor Moiseev, Rico Sennrich, and Ivan Titov. Analyzing Multi-Head Self-Attention: Specialized Heads Do the Heavy Lifting, the Rest Can Be Pruned. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019. URL <https://arxiv.org/abs/1905.09418>.
- [362] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, et al. Voyager: An Open-Ended Embodied Agent with Large Language Models. *arXiv Preprint arXiv:2305.16291*, 2023. URL <https://arxiv.org/abs/2305.16291>.
- [363] Junkang Wang, Jianfei Duan, Yue Liu, Zhangchen Yue, Hanghang Tong, and James Wang. Beyond Reverse KL: Generalizing Direct Preference Optimization with Diverse Divergence Constraints. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2309.16240>.

- [364] Lei Wang, Wanyu Xu, Yiber Lan, et al. Plan-and-Solve Prompting: Improving Zero-Shot Chain-of-Thought Reasoning by Large Language Models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023. URL <https://arxiv.org/abs/2305.04091>.
- [365] Peiyi Wang, Lei Li, Zhihong Shao, et al. Math-Shepherd: Verify and Reinforce LLMs Step-by-Step Without Human Annotations. *arXiv Preprint arXiv:2312.08935*, 2024. URL <https://arxiv.org/abs/2312.08935>.
- [366] Xingyao Wang, Boxuan Ding, Ziniu Hoang, et al. OpenHands: An Open Platform for AI Software Developers as Generalist Agents. *arXiv Preprint arXiv:2407.16741*, 2024.
- [367] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, et al. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *Proceedings of the 11th International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2203.11171>.
- [368] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, et al. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2203.11171>.
- [369] Yiduo Wang et al. RLFactory: A Plug-and-Play Framework for Multi-Round Tool-Use RL. *arXiv preprint*, 2025.
- [370] Yifan et al Wang. Improving Data Efficiency for LLM Reinforcement Fine-Tuning via Difficulty-Targeted Online Data Selection. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. URL <https://arxiv.org/abs/2506.05316>.
- [371] Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, et al. Self-Instruct: Aligning Language Models with Self-Generated Instructions. *arXiv Preprint arXiv:2212.10560*, 2022. URL <https://arxiv.org/abs/2212.10560>.
- [372] Zhihao Wang et al. Auxiliary-Loss-Free Load Balancing Strategy for Mixture-of-Experts. *arXiv preprint*, 2024.
- [373] Christopher J. C. H Watkins. Learning from Delayed Rewards. 1989.
- [374] Jason Wei, Xuezhi Wang, Dale Schuurmans, et al. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2201.11903>.
- [375] Gerhard Weiss and ed. Multiagent Systems: A Modern Approach to Distributed Artificial Intelligence, 1999.
- [376] Leandro von Werra, Younes Belkada, Lewis Tunstall, et al. TRL: Transformer Reinforcement Learning, 2022. URL <https://github.com/huggingface/trl>.
- [377] Jason Weston, Sumit Chopra, and Antoine Bordes. Memory Networks. In *International Conference on Learning Representations (ICLR)*, 2015. URL <https://arxiv.org/abs/1410.3916>.
- [378] Brandon T. Willard and Rémi Louf. Efficient Guided Generation for Large Language Models. *arXiv Preprint arXiv:2307.09702*, 2023. URL <https://arxiv.org/abs/2307.09702>.
- [379] Ronald J Williams. Simple Statistical Gradient-Following Algorithms for Connectionist Reinforcement Learning. *Machine Learning*, 1992. URL <https://link.springer.com/article/10.1007/BF00992696>.
- [380] Edwin B Wilson. Probable Inference, the Law of Succession, and Statistical Inference. *Journal of the American Statistical Association*, 1927. URL <https://www.jstor.org/stable/2276774>.

- [381] Michael Wooldridge. An Introduction to MultiAgent Systems, 2009.
- [382] Michael Wooldridge, Nicholas R. Jennings, and David Kinny. The Gaia Methodology for Agent-Oriented Analysis and Design. *Autonomous Agents and Multi-Agent Systems*, 2000.
- [383] BigScience Workshop. BLOOM: A 176B-Parameter Open-Access Multilingual Language Model. *arXiv Preprint arXiv:2211.05100*, 2023. URL <https://arxiv.org/abs/2211.05100>.
- [384] Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. LongMemEval: Benchmarking Chat Assistants on Long-Term Interactive Memory. In *International Conference on Learning Representations (ICLR)*, 2025. URL <https://arxiv.org/abs/2410.10813>.
- [385] Qingyun Wu, Gagan Bansal, Jieyu Zhang, et al. AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. *arXiv Preprint arXiv:2308.08155*, 2023.
- [386] Yifan Wu et al. Long-Horizon-Terminal-Bench: Dense Reward Grading for Command-Line Agents. *arXiv preprint*, 2026.
- [387] Zhengbao Wu, Aman Arora, Zhiqiang Wang, Byung-Gon Kim, and Tian Huang. Retrieval Head Mechanistically Explains Long-Context Factuality. *arXiv Preprint arXiv:2404.15574*, 2024. URL <https://arxiv.org/abs/2404.15574>.
- [388] Zhenyu Wu, Qinghua Hu, Yin Zhang, Yiming Gao, and Jiangtao Chen. An Empirical Analysis of Compute-Optimal Inference for Problem-Solving with Language Models. *arXiv Preprint arXiv:2408.00724*, 2024.
- [389] xAI. Grok 4.5: Training on Real Developer Session Traces. Technical report, 2026.
- [390] Zhiheng Xi, Yiwen Ding, Wenxiang Chen, et al. AgentGym: Evolving Large Language Model-Based Agents Across Diverse Environments. *arXiv Preprint arXiv:2406.04151*, 2024. URL <https://arxiv.org/abs/2406.04151>.
- [391] Fen Xia, Tie-Yan Liu, Jue Wang, Wensheng Zhang, and Hang Li. Listwise Approach to Learning to Rank: Theory and Algorithm. In *Proceedings of the 25th International Conference on Machine Learning (ICML)*, 2008.
- [392] Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2211.10438>.
- [393] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient Streaming Language Models with Attention Sinks. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2309.17453>.
- [394] Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient Streaming Language Models with Attention Sinks. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2309.17453>.
- [395] Shitao Xiao, Zheng Liu, Peitian Zhang, and Niklas Muennighoff. C-Pack: Packaged Resources to Advance General Chinese Embedding. *arXiv Preprint arXiv:2309.07597*, 2023.
- [396] Sang Michael Xie, Aditi Raghunathan, Percy Liang, and Tengyu Ma. An Explanation of in-Context Learning as Implicit Bayesian Inference. In *Proceedings of the 10th International Conference on Learning Representations (ICLR)*, 2022. URL <https://arxiv.org/abs/2111.02080>.
- [397] Tianbao Xie, Danyang Zhang, Jixuan Chen, et al. OSWorld: Benchmarking Multimodal Agents for Open-Ended Tasks in Real Computer Environments. In *Advances in Neural Information Processing Systems*, 2024.

- [398] Can Xu, Qingfeng Sun, Kai Zheng, et al. WizardLM: Empowering Large Language Models to Follow Complex Instructions. *arXiv Preprint arXiv:2304.12244*, 2023. URL <https://arxiv.org/abs/2304.12244>.
- [399] Haoxiang Xu, Hongyuan Zhao, Ying Liu, and Dong Wei. It Takes Two: Pairwise Preference Optimization with Two Rollouts. *arXiv Preprint arXiv:2505.07856*, 2025. URL <https://arxiv.org/abs/2505.07856>.
- [400] Wujiang Xu, Zujie Liang, Kai Mei, Hang Gao, Juntao Tan, and Yongfeng Zhang. A-Mem: Agentic Memory for LLM Agents. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- [401] Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. Corrective Retrieval Augmented Generation. *arXiv Preprint arXiv:2401.15884*, 2024. URL <https://arxiv.org/abs/2401.15884>.
- [402] Chengrun Yang, Xuezhi Wang, Yifeng Lu, et al. Large Language Models as Optimizers. In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2309.03409>.
- [403] Jingwen Yang, Yuxuan Zhu, Binyuan Wang, et al. ARQ: Attentive Reasoning Queries for Multi-Hop Question Answering over Long Contexts. *arXiv Preprint arXiv:2501.08290*, 2025. URL <https://arxiv.org/abs/2501.08290>.
- [404] John Yang, Carlos E. Jimenez, Alexander Wettig, et al. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL <https://arxiv.org/abs/2405.15793>.
- [405] Zhilin Yang, Peng Qi, Saizheng Zhang, et al. HotpotQA: A Dataset for Diverse, Explainable Multi-Hop Question Answering. In *Proceedings of EMNLP*, 2018.
- [406] Zhilin Yang, Zihang Dai, Yiming Yang, Jaime Carbonell, Ruslan Salakhutdinov, and Quoc V. Le. XLNet: Generalized Autoregressive Pretraining for Language Understanding. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [407] Shunyu Yao, Dian Yu, Jeffrey Zhao, et al. Tree of Thoughts: Deliberate Problem Solving with Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2305.10601>.
- [408] Shunyu Yao, Jeffrey Zhao, Dian Yu, et al. ReAct: Synergizing Reasoning and Acting in Language Models. In *International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2210.03629>.
- [409] Shunyu Yao, Dian Yu, Jeffrey Zhao, et al. Tree of Thoughts: Deliberate Problem Solving with Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL <https://arxiv.org/abs/2305.10601>.
- [410] Zhewei Yao, Samyam Rajbhandari, Reza Yazdani Aminabadi, et al. DeepSpeed-Chat: Easy, Fast and Affordable RLHF Training of ChatGPT-Like Models at All Scales. *arXiv Preprint arXiv:2308.01320*, 2023.
- [411] Jaeseok Yoo and Dongmin Shin. Adaptive Episode Length Adjustment for Multi-Agent Reinforcement Learning. In *Proceedings of the 24th International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, 2025.
- [412] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A Distributed Serving System for Transformer-Based Generative Models. In *Proceedings of the 16th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 2022.

- [413] Qiying Yu, Zheng Sun, Shang Wen, et al. DAPO: An Open-Source LLM Reinforcement Learning System. *arXiv Preprint arXiv:2503.14476*, 2025. URL <https://arxiv.org/abs/2503.14476>.
- [414] Xiao Yu, Baolin Peng, Ruize Xu, et al. Reinforcement World Model Learning for LLM-Based Agents. *arXiv Preprint arXiv:2602.05842*, 2026.
- [415] Ted Zadouri, Jay Shah, Ganesh Bikshandi, and Tri Dao. FlashAttention-4: Hardware-Efficient Attention on Blackwell GPUs with Minimal Software Design. *arXiv Preprint arXiv:2603.05451*, 2026. URL <https://arxiv.org/abs/2603.05451>.
- [416] Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, et al. Big Bird: Transformers for Longer Sequences. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. URL <https://arxiv.org/abs/2007.14062>.
- [417] Elad Ben Zaken, Shauli Ravfogel, and Yoav Goldberg. BitFit: Simple Parameter-Efficient Fine-Tuning for Transformer-Based Masked Language-Models. In *Proceedings of the 60th Annual Meeting of the ACL*, 2022. URL <https://arxiv.org/abs/2106.10199>.
- [418] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. STaR: Bootstrapping Reasoning with Reasoning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2203.14465>.
- [419] Eric Zelikman, Georges Harik, Yijia Shao, Varuna Jayasiri, Nick Haber, and Noah D. Goodman. Quiet-STaR: Language Models Can Teach Themselves to Think Before Speaking. *arXiv Preprint arXiv:2403.09629*, 2024. URL <https://arxiv.org/abs/2403.09629>.
- [420] Alex L. Zhang, Seyyed Hasan Mahdavi, Percy Liang, and Tatsunori Hashimoto. Recursive Language Models. *arXiv Preprint arXiv:2512.24601*, 2025. URL <https://arxiv.org/abs/2512.24601>.
- [421] Biao Zhang and Rico Sennrich. Root Mean Square Layer Normalization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019. URL <https://arxiv.org/abs/1910.07467>.
- [422] Jiayi Zhang, Simon Yu, Derek Chong, et al. Verbalized Sampling: How to Mitigate Mode Collapse and Unlock LLM Diversity. *arXiv Preprint arXiv:2510.01171*, 2025. URL <https://arxiv.org/abs/2510.01171>.
- [423] Kaiqing Zhang, Zhuoran Yang, and Tamer Başar. Multi-Agent Reinforcement Learning: A Selective Overview of Theories and Algorithms. *Handbook of Reinforcement Learning and Control*, 2021. URL <https://arxiv.org/abs/1911.10635>.
- [424] Qingru Zhang, Minshuo Chen, Alexander Bukharin, et al. AdaLoRA: Adaptive Budget Allocation for Parameter-Efficient Fine-Tuning. *arXiv Preprint arXiv:2303.10512*, 2023. URL <https://arxiv.org/abs/2303.10512>.
- [425] Tianjun Zhang, Shishir G. Patil, Naman Jain, et al. RAFT: Adapting Language Model to Domain Specific RAG. *arXiv Preprint arXiv:2403.10131*, 2024. URL <https://arxiv.org/abs/2403.10131>.
- [426] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. BERTScore: Evaluating Text Generation with BERT. In *International Conference on Learning Representations (ICLR)*, 2020. URL <https://arxiv.org/abs/1904.09675>.
- [427] Wei Zhang et al. Self-Evolving Software Agents: Combining BDI Architectures with LLMs. *arXiv preprint*, 2026.
- [428] Xinrong Zhang, Yingfa Chen, Shengding Hu, et al. InfiniteBench: Extending Long Context Evaluation Beyond 100K Tokens. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. URL <https://arxiv.org/abs/2402.13718>.

- [429] Zhenyu Zhang, Ying Sheng, Tianyi Zhou, et al. H2O: Heavy-Hitter Oracle for Efficient Generative Inference of Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2306.14048>.
- [430] Siyan Zhao, Zhihui Xie, Mengchen Liu, Jing Huang, Guan Pang, Feiyu Chen, and Aditya Grover. Self-Distilled Reasoner: On-Policy Self-Distillation for Large Language Models. In *Proceedings of ICLR 2026*, 2026. URL <https://arxiv.org/abs/2601.18734>. arXiv preprint arXiv:2601.18734.
- [431] Yanli Zhao, Andrew Gu, Rohan Varma, et al. PyTorch FSDP: Experiences on Scaling Fully Sharded Data Parallel. In *Proceedings of the VLDB Endowment*, 2023.
- [432] Yao Zhao, Rishabh Joshi, Tianqi Liu, Misha Khalman, Mohammad Saleh, and Peter J. Liu. SLiC-HF: Sequence Likelihood Calibration with Human Feedback. *arXiv Preprint arXiv:2305.10425*, 2023. URL <https://arxiv.org/abs/2305.10425>.
- [433] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2306.05685>.
- [434] Zhipu AI. GLM-4.5: Agentic, Reasoning, and Coding (ARC) Foundation Models. *arXiv preprint arXiv:2508.06471*, 2025. URL <https://arxiv.org/abs/2508.06471>.
- [435] Zhipu AI. GLM-5: From Vibe Coding to Agentic Engineering. *arXiv preprint arXiv:2602.15763*, 2026. URL <https://arxiv.org/abs/2602.15763>.
- [436] Yanqi Zhong, Yifu Chen, Zijun Li, and Minlie Chen. Importance Sampling Corrections for Large Language Model Alignment with Asynchronous Generation. *arXiv Preprint arXiv:2503.09057*, 2025. URL <https://arxiv.org/abs/2503.09057>.
- [437] Yanqi Zhong, Jiaqi Shi, Yifu Chen, and Minlie Chen. GDPO: Learning to Directly Align Language Models with Group-Decoupled Reward. *arXiv Preprint arXiv:2501.17888*, 2025. URL <https://arxiv.org/abs/2501.17888>.
- [438] Andy Zhou, Kai Yan, Michal Shlapentokh-Rothman, Haohan Wang, and Yu-Xiong Wang. Language Agent Tree Search Unifies Reasoning Acting and Planning in Language Models. *arXiv Preprint arXiv:2310.04406*, 2024.
- [439] Chunting Zhou, Pengfei Liu, Puxin Xu, et al. LIMA: Less Is More for Alignment. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. URL <https://arxiv.org/abs/2305.11206>.
- [440] Shuyan Zhou, Frank F. Xu, Hao Zhu, et al. WebArena: A Realistic Web Environment for Building Autonomous Agents. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2307.13854>.
- [441] Yongchao Zhou, Andrei Ioan Muresanu, Ziyen Han, et al. Large Language Models Are Human-Level Prompt Engineers. In *Proceedings of the 11th International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2211.01910>.
- [442] Daniel M. Ziegler, Nisan Stiennon, Jeffrey Wu, et al. Fine-Tuning Language Models from Human Preferences. *arXiv Preprint arXiv:1909.08593*, 2019. URL <https://arxiv.org/abs/1909.08593>.
- [443] Andy Zou, Zifan Wang, J. Zico Kolter, and Matt Fredrikson. Universal and Transferable Adversarial Attacks on Aligned Language Models. *arXiv Preprint arXiv:2307.15043*, 2023. URL <https://arxiv.org/abs/2307.15043>.